

Minsoo Kim



+8210-7714-9201



vibrio0102@gmail.com



[linkedin](#)



[github](#)

Summary

Data engineer who learned a new domain, interviewed researchers to define requirements, and designed and built an in-house MLOps platform from scratch. Independently learned the medical domain and on-premises MLOps, leading stack selection, hardware configuration and data-pipeline implementation. Previously built AWS data pipelines and led a data engineering team.

Selected Platform Work

Led requirements discovery and greenfield MLOps delivery: Interviewed researchers and examined data, code and execution workflows to identify producer dependencies, scattered datasets and missing versions and lineage. Assigned clear responsibilities to Airflow for execution and retries, dbt for SQL transformations, Iceberg for table versions and DataHub for lineage, then designed and implemented the shared platform.

Built and operated on-premises bare-metal Kubernetes: Sole owner of cluster construction on a two-person platform team. Built a 10-node cluster from company hardware and codified networking, storage, identity and execution environments in 12 Pulumi stacks (as of September 2026). Tracked infrastructure changes and inter-stack dependencies to support repeatable deployment.

Led hardware configuration and operations: Disassembled idle PCs and workstations to reclaim and reinstall RAM and NVMe drives; selected and replaced network equipment. Addressed NUMA and dual-channel configuration, memory pressure and swap, and separated the PostgreSQL data disk, adapting both hardware and system settings to actual workloads.

Designed Iceberg catalog and DuckDB query access: Used SeaweedFS S3-compatible storage and its Iceberg REST Catalog to manage Silver/Gold data and metadata with NVMe/HDD storage tiers. Built PostgreSQL + pg_duckdb query access. For dimension-table generation, resolved the latest metadata location from the catalog and passed it to DuckDB, avoiding version inference from filenames.

Earlier: Led AWS EMR/PySpark art-market ETL pipelines at Eazel as a team lead.

Current Focus

Data Engineer, HUIINNO (Dec 2025 -). Operate the shared data pipelines and on-premises platform while extending reproducibility features that connect data versions, code and environments. MLflow tracking, the model registry and manifest freeze/restore tools are built; research-code integration and end-to-end reproducibility validation are in progress.

Stack

Platform: Kubernetes (kubeadm) · Pulumi · Cilium · SeaweedFS · Dex/OIDC · Prometheus/Grafana

Data: Airflow 3.x · dbt · Apache Iceberg · DataHub · PostgreSQL (pg_duckdb) · ClickHouse

ML: Ray/KubeRay · MLflow

Languages & other: Python · SQL · JavaScript · Swift · Shell · AWS/GCP

Certifications & Languages

Certifications: Engineer Information Processing (Nov 2017) · Craftsman Information Processing (Jul 2009) — issued by HRD Korea

Languages: Korean (native) · English OPIc AL · Japanese (conversational)

Full experience and project details are in a separate document ([resume_english.pdf](#)).